

04_deployment_garch

May 15, 2026

1 04 — GARCH(1,1)-t Deployment Model

The benchmark (notebooks 01–03) concluded that the nine top models are statistically indistinguishable on *point* forecasts. So how do we choose one for deployment?

The tiebreaker is prediction intervals. Simple baselines (Naïve, RWD, ETS, Theta) produce point forecasts only, or intervals that assume constant variance and therefore ignore volatility regimes. GARCH-family models give *calibrated* prediction intervals that adapt to market stress. For a forecast intended for practical use — hedging or procurement planning — interval calibration matters more than a small difference in point MAE.

This notebook validates each modeling choice with a statistical test:

Decision	Justified by
Model log-returns, not prices	ADF test (already shown in 03_stationarity_diagnostics.ipynb)
Constant mean, not AR	HAC-robust t-stat on the drift
Student-t innovations, not Normal	Excess kurtosis ~ 6.6; Jarque-Bera
GARCH family over IID	ARCH LM test; ACF of squared returns
GARCH(1,1)-t specification	Standard symmetric volatility model with heavy-tailed innovations
Expanding-window backtest	60-origin empirical coverage check

1.1 1. Setup

```
[ ]: import sys
      from pathlib import Path

      # Make `coffee_forecast` importable whether or not `pip install -e .`
      # is in effect for this kernel. No-op when the package is already on
      # sys.path (e.g. via the editable install); falls back to inserting
      # the repo root so the notebook also runs from a fresh clone.
      try:
          import coffee_forecast # noqa: F401
      except ModuleNotFoundError:
          sys.path.insert(0, str(Path.cwd().parent))

      import warnings
```

```
warnings.filterwarnings("ignore", category=FutureWarning)

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from arch import arch_model
from scipy import stats
from statsmodels.stats.diagnostic import acorr_ljungbox, het_arch

from coffee_forecast import load_coffee_data, forecast, std_t_ppf, plot_forecast
from coffee_forecast.config import (
    CSV_DIR,
    FIG_DIR,
    HORIZON,
    CONTEXT_LEN,
    REPO_ROOT,
)

CSV_DIR.mkdir(parents=True, exist_ok=True)
FIG_DIR.mkdir(parents=True, exist_ok=True)
```

```
[2]: df = load_coffee_data()
prices = df["y"].values.astype(float)
log_returns = np.log(df["y"]).diff().dropna().values
print(f"{len(prices)} prices  ({df['ds'].min().date()} → {df['ds'].max().
    ↳date()})")
print(f"{len(log_returns)} daily log-returns")
```

```
8104 prices  (1994-01-03 → 2026-01-30)
8103 daily log-returns
```

1.2 2. Drift test — does coffee have a statistically significant trend?

Coffee prices have risen substantially over 32 years, but *daily* returns average only +0.019%/day — around 4.7% annualized, which sounds like real drift. To check whether that's actually significant we use HAC-robust standard errors (Newey-West), which correct for any autocorrelation or heteroskedasticity in the return series.

```
[3]: import statsmodels.api as sm

y = log_returns
X = np.ones_like(y) # regression on a constant = mean
res_hac = sm.OLS(y, X).fit(cov_type="HAC", cov_kwds={"maxlags": 10})

mu, se = res_hac.params[0], res_hac.bse[0]
t_stat, p_val = mu / se, res_hac.pvalues[0]
print(f"Mean daily log-return:  {mu:+.6f}")
```

```

print(f"HAC standard error:      {se:.6f}")
print(f"t-stat:                  {t_stat:+.3f}")
print(f"p-value (HAC):           {p_val:.3f}")
verdict = "not significant" if p_val > 0.05 else "significant"
print(f"\n→ Drift is statistically {verdict} at 5%. Use a zero-mean model.")

```

```

Mean daily log-return:    +0.000188
HAC standard error:       0.000255
t-stat:                   +0.737
p-value (HAC):            0.461

```

→ Drift is statistically not significant at 5%. Use a zero-mean model.

1.3 3. Distribution fit — Normal vs. Student-t

Daily log-returns famously have fat tails. Before choosing a GARCH innovation distribution, we quantify just how fat they are.

```

[4]: skew = stats.skew(log_returns)
kurt = stats.kurtosis(log_returns, fisher=True) # 0 under Normal
jb, jb_p = stats.jarque_bera(log_returns)
print(f"Skewness:      {skew:+.3f}")
print(f"Excess kurtosis: {kurt:+.3f} (Normal = 0)")
print(f"Jarque-Bera stat: {jb:.1f} (p = {jb_p:.2e})")
print(f"\n→ Excess kurtosis of ~6.6 is wildly non-Normal. Use Student-t.")

```

```

Skewness:      +0.186
Excess kurtosis: +6.648 (Normal = 0)
Jarque-Bera stat: 14970.2 (p = 0.00e+00)

```

→ Excess kurtosis of ~6.6 is wildly non-Normal. Use Student-t.

```

[5]: nu_fit, t_loc, t_scale = stats.t.fit(log_returns)
mu_norm, std_norm = stats.norm.fit(log_returns)

fig, axes = plt.subplots(1, 2, figsize=(13, 4.5), dpi=180)

# Panel (a): histogram with Normal and Student-t overlays
ax = axes[0]
x_grid = np.linspace(log_returns.min(), log_returns.max(), 500)
ax.hist(
    log_returns,
    bins=150,
    density=True,
    alpha=0.5,
    color="steelblue",
    label="Observed log-returns",
)

```

```


ax.plot(
    x_grid,
    stats.norm.pdf(x_grid, mu_norm, std_norm),
    "r--",
    lw=2,
    label=f"Normal fit (\u03bc={mu_norm:.4f}, \u03c3={std_norm:.4f})",
)
ax.plot(
    x_grid,
    stats.t.pdf(x_grid, nu_fit, t_loc, t_scale),
    color="#E67E22",
    lw=2,
    label=f"Student-t fit (\u03bd={nu_fit:.2f})",
)
ax.axvline(0, color="gray", ls="--", alpha=0.4)
ax.set_title("(a) Histogram vs. Normal vs. Student-t", fontsize=11,
    weight="semibold")
ax.set_xlabel("Log return")
ax.set_ylabel("Density")
ax.legend(fontsize=8, frameon=False)
ax.grid(alpha=0.2)
ax.spines["top"].set_visible(False)
ax.spines["right"].set_visible(False)

# Panel (b): Q-Q vs Normal
ax = axes[1]
(osm, osr), (slope, intercept, _) = stats.probplot(log_returns, dist="norm")
ax.scatter(osm, osr, s=2, color="steelblue", alpha=0.5)
ax.plot(osm, slope * np.array(osm) + intercept, "r--", lw=1.5, label="Normal
    reference")
ax.set_title("(b) Q-Q plot vs. Normal", fontsize=11, weight="semibold")
ax.set_xlabel("Theoretical quantiles")
ax.set_ylabel("Sample quantiles")
ax.legend(fontsize=8, frameon=False)
ax.grid(alpha=0.2)
ax.spines["top"].set_visible(False)
ax.spines["right"].set_visible(False)

fig.tight_layout()
fig.savefig(FIG_DIR / "returns_distribution.png", dpi=300, bbox_inches="tight")
plt.show()

print(
    f"Fitted Student-t \u03bd = {nu_fit:.2f} \u2192 tails substantially
    heavier than Normal."
)

```



Fitted Student-t = 4.21 → tails substantially heavier than Normal.

The histogram shows the data's peak is higher than Normal and the tails are fatter; the Q-Q plot's S-shape at the extremes is the classic fingerprint of heavy tails. The Student-t fit with 4.2 absorbs this tail mass in a single parameter. This is what motivates Student-t innovations in the GARCH specification below.

1.4 4. ARCH test — is there volatility clustering?

Engle's (1982) ARCH LM test checks whether squared returns have detectable autocorrelation. A rejection at 5% says there's detectable volatility clustering — a necessary condition for GARCH to be better than an IID model.

```
[6]: lm_stat, lm_p, _, _ = het_arch(log_returns, nlags=10)
print(f"ARCH LM test (10 lags): statistic = {lm_stat:.1f}, p = {lm_p:.2e}")
print(
    "\n→ Strongly reject the null of no volatility clustering. GARCH family is_
    appropriate."
)
```

ARCH LM test (10 lags): statistic = 523.2, p = 4.96e-106

→ Strongly reject the null of no volatility clustering. GARCH family is appropriate.

```
[7]: from statsmodels.graphics.tsaplots import plot_acf

fig, axes = plt.subplots(1, 2, figsize=(13, 4), dpi=180)

plot_acf(
    log_returns,
    lags=40,
    ax=axes[0],
    alpha=0.05,
    zero=False,
```

```


    title="(a) ACF of log-returns",
    color="steelblue",
)
axes[0].set_xlabel("Lag (trading days)")
axes[0].set_ylim(-0.1, 0.25)
axes[0].grid(alpha=0.2)

plot_acf(
    log_returns**2,
    lags=40,
    ax=axes[1],
    alpha=0.05,
    zero=False,
    title="(b) ACF of squared log-returns",
    color="#E67E22",
)
axes[1].set_xlabel("Lag (trading days)")
axes[1].set_ylim(-0.1, 0.25)
axes[1].grid(alpha=0.2)

for ax in axes:
    ax.spines["top"].set_visible(False)
    ax.spines["right"].set_visible(False)

fig.tight_layout()
fig.savefig(FIG_DIR / "acf_returns_vs_squared.png", dpi=300,
           bbox_inches="tight")
plt.show()

```



The contrast makes the modeling strategy concrete. Returns themselves are nearly white noise — the ACF bars mostly fall inside the confidence band — so a mean model has little to exploit. Squared returns, by contrast, show persistent positive autocorrelation out to dozens of lags: today’s magnitude predicts tomorrow’s magnitude, even though today’s direction doesn’t predict tomorrow’s direction. This is volatility clustering, and it is what GARCH is designed to capture.

1.5 5. Fit GARCH(1,1)-t

The symmetric GARCH(1,1) with Student-t innovations:

$$h_t = \omega + \alpha \varepsilon_{t-1}^2 + \beta h_{t-1}, \quad z_t \sim t_\nu$$

Volatility depends on the last squared shock plus the last variance level, weighted equally regardless of shock sign. Combined with Student-t innovations to handle the heavy tails documented in §3, this is the deployment specification.

```
[8]: # arch_model needs returns in percent scale for numerical stability.
r_pct = log_returns * 100.0

am_garch = arch_model(r_pct, mean="Zero", vol="Garch", p=1, q=1, dist="t")
garch_fit = am_garch.fit(disp="off")
print(garch_fit.summary())
```

Zero Mean - GARCH Model Results

```
=====
=====
Dep. Variable:                y    R-squared:
0.000
Mean Model:                Zero Mean    Adj. R-squared:
0.000
Vol Model:                GARCH    Log-Likelihood:
-17741.5
Distribution:    Standardized Student's t    AIC:
35490.9
Method:                Maximum Likelihood    BIC:
35518.9
                                No. Observations:
8103
Date:                Thu, May 14 2026    Df Residuals:
8103
Time:                22:46:21    Df Model:
0
```

Volatility Model

```
=====
=====
              coef    std err          t      P>|t|      95.0% Conf. Int.
-----
omega          0.0992   3.742e-02     2.651   8.036e-03   [2.584e-02,  0.173]
alpha[1]       0.0409   8.729e-03     4.681   2.854e-06   [2.375e-02, 5.797e-02]
beta[1]        0.9413   1.463e-02    64.335   0.000         [ 0.913,  0.970]
```

Distribution

```
=====
=====
              coef    std err          t      P>|t|      95.0% Conf. Int.
-----
nu            5.5075     0.342     16.086   3.215e-58   [ 4.836,  6.179]
```

=====

Covariance estimator: robust

1.6 6. Residual diagnostics

The final sanity check. If the model fits, standardized residuals ε_t/σ_t should look roughly like iid Student-t draws.

```
[9]: alpha_c = garch_fit.params["alpha[1]"]
     beta_c = garch_fit.params["beta[1]"]
     nu_hat = garch_fit.params["nu"]
     persistence = alpha_c + beta_c

     lb = acorr_ljungbox(garch_fit.std_resid**2, lags=[10, 20], return_df=True)

     print(f"Volatility persistence ( + ):           {persistence:.4f}")
     print(
         f"Half-life of volatility shocks:           {np.log(2) / -np.log(persistence):.
         ↪0f} days"
     )
     print(f"Fitted (tail heaviness):                 {nu_hat:.2f}")
     print(f"Ljung-Box on squared residuals:")
     print(f"  LB(10) p-value:                         {lb['lb_pvalue'].iloc[0]:.4f}")
     print(f"  LB(20) p-value:                         {lb['lb_pvalue'].iloc[1]:.4f}")
```

```
Volatility persistence ( + ):           0.9822
Half-life of volatility shocks:          39 days
Fitted (tail heaviness):                 5.51
Ljung-Box on squared residuals:
  LB(10) p-value:                        0.0003
  LB(20) p-value:                        0.0009
```

Reading the diagnostics.

- **Persistence ~0.98** (half-life on the order of weeks) is typical for daily commodity volatility — shocks fade slowly but the process is stable (persistence < 1 is required for the unconditional variance to exist).
- The fitted $\nu \approx 5.5$ is *higher* than §3's standalone Student-t fit on raw returns ($\nu \approx 4.2$). They measure different quantities: §3's ν is the tail thickness of *unconditional* daily returns; this ν is the tail thickness of *standardized residuals* after GARCH filters out the time-varying volatility. Heavy unconditional tails come partly from volatility clustering itself, and once GARCH absorbs that the residual distribution looks closer to Normal — though still heavy enough that Student-t innovations are required.
- The Ljung-Box p-values on squared residuals are the cleanest summary of whether the conditional-variance specification has absorbed the volatility-clustering structure documented in §4.

1.7 7. Conditional vs. constant-variance intervals

The residual diagnostics above show the model fits the data. The next question is what that fit buys us in practice. The plot below contrasts a fixed $\pm 1.96 \cdot \hat{\sigma}$ band (what a constant-variance model would produce) with the GARCH(1,1)-t adaptive $\pm 1.96 \cdot \sigma_t$ band. Both hit $\approx 95\%$ unconditional coverage; what differs is *when* they're wide and *when* they're narrow.

```
[ ]: # Work in % scale to match the fitted model.
r_pct = log_returns * 100
cond_sd = garch_fit.conditional_volatility
uncond_sd = r_pct.std()

# Both bands use the textbook normal 95% quantile (1.96). The contrast
# is constant vs. conditional sigma, holding the quantile fixed.
Z95 = 1.96

const_upper = +Z95 * uncond_sd
const_lower = -Z95 * uncond_sd
cond_upper = +Z95 * cond_sd
cond_lower = -Z95 * cond_sd

dates = df["ds"].iloc[1:].values

fig, ax = plt.subplots(figsize=(13, 5), dpi=180)
ax.plot(dates, r_pct, color="gray", lw=0.4, alpha=0.7, label="Daily log-return")
ax.axhline(
    const_upper,
    color="crimson",
    ls="--",
    lw=1.5,
    label="Constant-variance  $\pm 1.96 \cdot \hat{\sigma}$ ",
)
ax.axhline(const_lower, color="crimson", ls="--", lw=1.5)
ax.plot(
    dates,
    cond_upper,
    color="teal",
    lw=1.2,
    label="GARCH(1,1)-t  $\pm 1.96 \cdot \sigma_t$ ",
)
ax.plot(dates, cond_lower, color="teal", lw=1.2)
ax.axhline(0, color="black", lw=0.5, alpha=0.4)


ax.set_title(
    "Constant-variance vs. conditional 95% bounds on daily returns",
    fontsize=12,
    weight="semibold",
)
```

```

ax.set_xlabel("Date")
ax.set_ylabel("Daily log-return (%)")
ax.legend(loc="upper right", fontsize=9, frameon=False)
ax.grid(alpha=0.2)
ax.spines["top"].set_visible(False)
ax.spines["right"].set_visible(False)
fig.tight_layout()
fig.savefig(
    FIG_DIR / "constant_vs_conditional_bounds.png", dpi=300, bbox_inches="tight"
)
plt.show()

inside_const = ((r_pct >= const_lower) & (r_pct <= const_upper)).mean()
inside_cond = ((r_pct >= cond_lower) & (r_pct <= cond_upper)).mean()
print(f"In-sample coverage at nominal 95%:")
print(f"  Constant-variance band:      {inside_const:.1%}")
print(f"  GARCH(1,1)-t conditional band: {inside_cond:.1%}")
print(f"_t range: {cond_sd.min():.2f}% - {cond_sd.max():.2f}% per day")
print(f"→ peak _t is {cond_sd.max() / uncond_sd:.1f}× the unconditional_
    ↪level")

```



```

In-sample coverage at nominal 95%:
  Constant-variance band:      94.8%
  GARCH(1,1)-t conditional band: 94.5%
_t range: 1.37% - 6.74% per day
→ peak _t is 2.8× the unconditional level

```

Both bands achieve roughly 95% unconditional coverage — a “broken clock is right twice a day” result. The important difference is *when* the coverage is achieved. The constant band is too wide in calm regimes (wasting precision when volatility is low) and too narrow in stressed regimes (failing to cover when volatility is high). The conditional band adapts, expanding during the 1997 Brazilian frost, the 2011 spike, and the 2024-2025 rally, and narrowing during the long 2002-2005 lull. For downstream users who

need to reason about tail risk in real time, this distinction is the difference between useful and misleading intervals.

1.8 8. Historical backtest — empirical coverage at 60 origins

Everything above justifies the *form* of the model. The ultimate test is whether its prediction intervals are *calibrated*: does a stated 95% PI actually contain 95% of realized observations?

We refit GARCH(1,1)-t at 60 evenly-spaced origins using **expanding** windows (all history up to each origin), forecast 63 business days, and check whether the actual path falls inside the 80% and 95% PIs.

Why expanding, not rolling: GARCH tail parameters (especially ν) are notoriously sample-sensitive. Short rolling windows give wildly varying ν across origins; expanding windows stabilize it, which empirically produces better-calibrated intervals.

Expect this cell to take a couple minutes — 60 MLE fits on 30+ years of data.

```
[11]: from coffee_forecast import get_forecast_origins

origins = get_forecast_origins(
    df, num_windows=60, context_len=CONTEXT_LEN, horizon=HORIZON
)
print(f"Fitting GARCH(1,1)-t at {len(origins)} origins...")

ALPHA_LEVELS = [0.50, 0.20, 0.10, 0.05, 0.01]  # 50%, 80%, 90%, 95%, 99% PIs
LEVELS_PCT = [int(round((1 - a) * 100)) for a in ALPHA_LEVELS]
coverage = {pct: [] for pct in LEVELS_PCT}
nus = []

for i, a in enumerate(origins, start=1):
    ctx_prices = df.iloc[:a]["y"].values.astype(float)
    ctx_log_ret = np.diff(np.log(ctx_prices))
    actual = df.iloc[a : a + HORIZON]["y"].values.astype(float)
    P0 = ctx_prices[-1]

    am = arch_model(ctx_log_ret * 100, mean="Zero", vol="Garch", p=1, q=1,
    ↪dist="t")
    fit = am.fit(disp="off", show_warning=False)
    nu_local = float(fit.params["nu"])
    nus.append(nu_local)

    var_fc = fit.forecast(horizon=HORIZON).variance.iloc[-1].values / 10_000.0
    cum_sd = np.sqrt(np.cumsum(var_fc))
    # Zero-mean spec: no drift term - point forecast is flat at P0.

    for alpha in ALPHA_LEVELS:
        q = std_t_ppf(1 - alpha / 2, nu_local)
        spread = q * cum_sd
```

```

        upper = P0 * np.exp(+spread)
        lower = P0 * np.exp(-spread)
        pct = int(round((1 - alpha) * 100))
        coverage[pct].append((actual >= lower) & (actual <= upper))

    if i % 10 == 0:
        print(f"    ... {i}/{len(origins)} origins done")

print()
print(
    f"GARCH(1,1)-t : min={min(nus):.2f} median={np.median(nus):.2f} "
    f"max={max(nus):.2f} std={np.std(nus):.2f}"
)

```

Fitting GARCH(1,1)-t at 60 origins...

```

... 10/60 origins done
... 20/60 origins done
... 30/60 origins done
... 40/60 origins done
... 50/60 origins done
... 60/60 origins done

```

GARCH(1,1)-t : min=3.99 median=4.63 max=5.50 std=0.45

```

[12]: fig, ax = plt.subplots(figsize=(8, 4.5), dpi=180)

print(f"{'PI':>5} {'Empirical':>11} {'Deviation':>11}")
print("-" * 31)
for pct in LEVELS_PCT:
    cov_matrix = np.array(coverage[pct])
    overall = cov_matrix.mean()
    deviation = overall - pct / 100
    print(f"{pct:>4}% {overall:>10.2%} {deviation:>10.2%}")

# Coverage-by-horizon figure: 80% and 95% only, to match the headline
# levels in S9's fan-chart panels. The other levels are summarized in
# the pooled table above.
for pct, color in [(80, "steelblue"), (95, "darkorange")]:
    cov_matrix = np.array(coverage[pct])
    by_step = cov_matrix.mean(axis=0)
    ax.plot(
        np.arange(1, HORIZON + 1),
        by_step * 100,
        color=color,
        lw=1.5,
        label=f"{pct}% observed",
    )

```

```


ax.axhline(pct, color=color, ls="--", alpha=0.5, label=f"{pct}% target")

ax.set_xlabel("Forecast horizon (days)")
ax.set_ylabel("Empirical coverage (%)")
ax.set_title(
    "GARCH(1,1)-t empirical coverage by horizon\nexpanding-window backtest at 60 origins",
    fontsize=11,
    weight="semibold",
)
ax.set_ylim(60, 105)
ax.grid(alpha=0.3)
ax.spines["top"].set_visible(False)
ax.spines["right"].set_visible(False)
ax.legend(loc="lower right", fontsize=8)

fig.tight_layout()
fig.savefig(FIG_DIR / "garch_coverage.png", dpi=300, bbox_inches="tight")
plt.show()

```

PI	Empirical	Deviation
50%	49.26%	-0.74%
80%	79.50%	-0.50%
90%	91.03%	+1.03%
95%	97.30%	+2.30%
99%	99.84%	+0.84%



Reading the coverage plot. The observed-coverage curves should hover near their targets (dashed lines). If the 95% line systematically sits below 95%, the model is *overconfident* on this data and should not be deployed. Coverage close to nominal across the full 63-day horizon is what licenses GARCH(1,1)-t as the deployment specification.

1.9 9. Sample forecasts across historical origins

A visual diagnostic of the backtest. Six origins are sampled at random from the 60 (seeded for reproducibility); for each we overlay the point forecast, 80% PI, 95% PI, and realized path.

```
[ ]: rng = np.random.default_rng(42)
sample_idx = sorted(rng.choice(len(origins), size=6, replace=False))

sample = []
for idx in sample_idx:
    a = origins[idx]
    ctx_prices = df.iloc[:a]["y"].values.astype(float)
    ctx_log_ret = np.diff(np.log(ctx_prices))
    actual = df.iloc[a : a + HORIZON]["y"].values.astype(float)
    origin_date = df.iloc[a]["ds"]

    am = arch_model(ctx_log_ret * 100, mean="Zero", vol="Garch", p=1, q=1,
    ↪dist="t")
    fit = am.fit(dis="off", show_warning=False)
    nu_local = float(fit.params["nu"])
    var_fc = fit.forecast(horizon=HORIZON).variance.iloc[-1].values / 10_000.0
    cum_sd = np.sqrt(np.cumsum(var_fc))
    P0 = ctx_prices[-1]
    # Zero-mean spec: point forecast is flat at P0.
    point = np.full(HORIZON, P0)

    q80 = std_t_ppf(0.90, nu_local)
    q95 = std_t_ppf(0.975, nu_local)
    sample.append(
        {
            "origin_id": idx + 1,
            "origin_date": origin_date,
            "actual": actual,
            "point": point,
            "lower_80": P0 * np.exp(-q80 * cum_sd),
            "upper_80": P0 * np.exp(+q80 * cum_sd),
            "lower_95": P0 * np.exp(-q95 * cum_sd),
            "upper_95": P0 * np.exp(+q95 * cum_sd),
        }
    )

all_prices = np.concatenate(
```

```

        [np.concatenate([s["actual"], s["upper_95"], s["lower_95"]]) for s in sample]
    )
    y_min, y_max = all_prices.min() * 0.95, all_prices.max() * 1.05

    fig, axes = plt.subplots(2, 3, figsize=(16, 8), dpi=180, sharex=True)
    steps = np.arange(1, HORIZON + 1)

    for ax, s in zip(axes.flat, sample):
        ax.fill_between(
            steps,
            s["lower_80"],
            s["upper_80"],
            color="steelblue",
            alpha=0.3,
            label="80% PI",
        )
        ax.fill_between(
            steps,
            s["lower_95"],
            s["upper_95"],
            color="steelblue",
            alpha=0.12,
            label="95% PI",
        )
        ax.plot(
            steps, s["point"], color="steelblue", ls="--", lw=1.5, label="Point_
            ↪forecast"
        )
        ax.plot(steps, s["actual"], color="black", lw=1.5, label="Actual")
        date_str = s["origin_date"].strftime("%b %Y")
        ax.set_title(f"Origin {s['origin_id']} - {date_str}", fontsize=10)
        ax.set_ylim(y_min, y_max)
        ax.grid(alpha=0.2)
        ax.spines["top"].set_visible(False)
        ax.spines["right"].set_visible(False)

    for ax in axes[-1]:
        ax.set_xlabel("Days ahead")
    for ax in axes[:, 0]:
        ax.set_ylabel("Price (cents/lb)")

    handles, labels = axes[0, 0].get_legend_handles_labels()
    fig.legend(
        handles,
        labels,
        loc="lower center",


```

```

ncol=4,
fontsize=10,
bbox_to_anchor=(0.5, -0.02),
frameon=False,
)

fig.suptitle(
    "GARCH(1,1)-t sample forecasts across six historical origins",
    fontsize=12,
    weight="semibold",
    y=1.00,
)
fig.tight_layout()
fig.savefig(FIG_DIR / "garch_sample_forecasts.png", dpi=300,
            bbox_inches="tight")
fig.savefig(FIG_DIR / "garch_sample_forecasts.pdf", bbox_inches="tight")
plt.show()

```



1.10 10. Sample forecast on committed history

With the model specification validated and coverage empirically confirmed, this section produces a *sample* forecast on the committed history alone, for documentation and reproducibility. Re-running the notebook always reproduces the same numbers, since `coffee_forecast.forecast(...)` is deterministic given the same input series.

This is **not** the live deployment output. The production forecast at `forecasts/latest_forecast.{csv,png}` (the image at the top of the README) is written exclusively by `scripts/run_forecast.py` running on a daily GitHub Actions schedule, which

calls `fetch_latest()` to pull newer Coffee C closes from Yahoo Finance (`KC=F`) before forecasting. The sample below uses only the prices already in `data/coffee.csv`, so its `run_date` will track that file rather than the current calendar day.

The sample is saved to `results/csv/garch_sample_forecast.csv` and `results/figures/garch_sample_forecast.{png,pdf}` — safely separate from the production paths so re-running this notebook never clobbers the live deployment artifacts.

```
[14]: fc = forecast(df, horizon=HORIZON)

print(f"Run date:      {fc['run_date'].iloc[0]}")
print(f"Last close:    {df['y'].iloc[-1]:.2f} cents/lb")
print(f"Day 1 point:    {fc['point'].iloc[0]:.2f}")
print(f"Day 63 point:   {fc['point'].iloc[-1]:.2f}")
print(f"Day 1 80% PI: [{fc['lo_80'].iloc[0]:.2f}, {fc['hi_80'].iloc[0]:.2f}]")
print(f"Day 63 95% PI: [{fc['lo_95'].iloc[-1]:.2f}, {fc['hi_95'].iloc[-1]:.2f}]\n")
fc.head()
```

```
Run date:      2026-01-30
Last close:    332.25 cents/lb
Day 1 point:   332.25
Day 63 point:  332.25
Day 1 80% PI:  [323.53, 341.21]
Day 63 95% PI: [230.11, 479.73]
```

```
[14]:
```

	run_date	target_date	horizon_days	point	ann_vol	lo_50	hi_50	\
0	2026-01-30	2026-02-02	1	332.25	36.36	327.90	336.66	
1	2026-01-30	2026-02-03	2	332.25	36.38	326.11	338.51	
2	2026-01-30	2026-02-04	3	332.25	36.40	324.74	339.93	
3	2026-01-30	2026-02-05	4	332.25	36.42	323.59	341.14	
4	2026-01-30	2026-02-06	5	332.25	36.44	322.58	342.21	

	lo_80	hi_80	lo_95	hi_95
0	323.53	341.21	317.40	347.79
1	319.98	344.99	311.44	354.45
2	317.28	347.93	306.94	359.64
3	315.02	350.43	303.20	364.09
4	313.04	352.64	299.93	368.05

```
[15]: # Sample-only paths: kept under results/, NOT forecasts/, so re-running
# this notebook never overwrites the production live-deployment files.
csv_path = CSV_DIR / "garch_sample_forecast.csv"
png_path = FIG_DIR / "garch_sample_forecast.png"
pdf_path = FIG_DIR / "garch_sample_forecast.pdf"
fc.to_csv(csv_path, index=False)
plot_forecast(df, fc, png_path)
plot_forecast(df, fc, pdf_path)
```

```


print(f"Wrote: {csv_path.relative_to(REPO_ROOT)}")
print(f"Wrote: {png_path.relative_to(REPO_ROOT)}")
print(f"Wrote: {pdf_path.relative_to(REPO_ROOT)}")

# Display the saved plot inline.
from IPython.display import Image, display

display(Image(filename=str(png_path)))

```

Wrote: results/csv/garch_sample_forecast.csv
Wrote: results/figures/garch_sample_forecast.png
Wrote: results/figures/garch_sample_forecast.pdf



1.11 Summary

Decision	Verdict
Work in log-return space	Confirmed (stationarity in 03_stationarity_diagnostics.ipynb)
Constant-mean specification	Drift not significant under HAC SEs
Student-t innovations	Excess kurtosis ~6.6 rules out Normal
GARCH over IID	ARCH LM rejects IID at any reasonable level
GARCH(1,1)-t specification	Standard symmetric model, calibrated under backtest
Expanding-window fitting	Stabilizes ν ; better coverage in backtest
80% / 95% PIs	Empirical coverage at or slightly above nominal

Deployment pipeline, in order: load committed history, append any newer closes from Yahoo,

fit GARCH(1,1)-t on log-returns, produce a 63-day forecast with 50/80/95% prediction intervals, write the CSV and the plot to `forecasts/`.

README image: whatever `scripts/run_forecast.py` produced on its last scheduled run — normally less than 24 hours old.